

Mime: Universal Speech, Visualized in Real-Time

COMP4107 Project Report

Winter 2026 - April 10, 2026
Carleton University

Group 10 - Daniel Lu, Methira Herath
101304181, 101304349
daniellu@cmail.carleton.ca, methiraherath@cmail.carleton.ca

Instructor: Matthew Holden

Table of Contents

Contributions	iii
1 Introduction	1
1.1 Background	1
1.2 Motivation and Objectives	1
1.3 Related Prior Work	2
2 Methods	3
2.1 Methods from Neural Networks	3
2.1.1 Speech Translation Pipeline	3
2.1.2 Audio-to-Blendshape Neural Model	3
2.1.3 Real-Time Inference and Rendering	4
2.2 Dataset and Data Processing	4
2.2.1 Primary Dataset	4
2.2.2 Preprocessing Pipeline	4
2.3 Validation Strategy and Experiments	5
2.3.1 Model-Level Validation	5
2.3.2 Hyperparameter Experiments	5
2.3.3 System-Level Validation	5
3 Results	7
3.1 Quantitative Data	7
3.2 Qualitative Results	7
4 Discussion	8
4.1 Limitations and Challenges	8
4.2 Directions for Future Work	8
4.3 Implications and Impact	9
4.4 Conclusions	9
References	10

Contributions

THIS IS A TEMPLATE SECTION, PLEASE REMIND MYSELF TO REPLACE LATER.

Chapter 1

Introduction

Language barriers in video communication are no longer only a text or audio problem. Recent progress in speech-to-text, machine translation, and text-to-speech has made multilingual conversations more accessible, but live conversations over video still feel unnatural when the spoken audio and visible mouth motion do not match. This mismatch introduces an *immersion gap*: listeners hear translated speech while seeing lip movements that correspond to a different language and timing. In practical settings such as online meetings, this visual-audio inconsistency can reduce comprehension, increase cognitive load, and make interactions feel less human.

This report presents **Mime**, a real-time multimodal communication system that chains speech translation with synchronized facial animation, running on consumer hardware with latency low enough for live calls. Rather than generating synthetic face frames from scratch, Mime predicts ARKit facial blendshape coefficients directly from audio features and renders them on an animated avatar: an approach that intentionally trades some visual fidelity for the speed that *live operation* requires.

1.1 Background

Human communication is strongly multimodal. Timing, facial movement, and visual articulation all contribute to how spoken content is perceived, and disruptions to that coherence are not merely aesthetic - they measurably affect comprehension and naturalness [1]. In cross-lingual video calls, even accurate translation leaves this visual layer broken: the listener hears one language while watching lip movements shaped by another.

State-of-the-art visual dubbing systems address this problem, but they were designed for film post-production. They assume offline workflows, high-end compute, and closed pipelines - none of which are compatible with live telepresence. Mime targets the practical gap: continuous microphone capture, real-time translation, audio-driven facial animation, and direct integration with standard conferencing stacks on Windows.

1.2 Motivation and Objectives

The project is motivated by a simple goal: make multilingual video conversation feel coherent in both sound and appearance. Instead of treating translation and lip-sync as separate tools, Mime integrates them into one stream-oriented system. The platform is designed around three core objectives:

-
1. **End-to-end integration:** Build a real-time pipeline that chains speech recognition, machine translation, and synthesized speech with an audio-to-blendshape visual model.
 2. **Temporal visual coherence:** Produce stable lip-sync by mapping speech features to 52 ARKit facial blendshape coefficients and rendering an animated avatar suitable for live overlay or virtual camera output.
 3. **Deployment practicality:** Maintain low latency on consumer-grade hardware, with integration support for common video conferencing workflows.

In implementation terms, Mime uses modular processors for STT, MT, and TTS, then routes synthesized audio to both playback and a fast inference path for facial animation. This architecture enables synchronized media output and reduces duplicated processing in live operation.

1.3 Related Prior Work

The main tools people have built for this are visual dubbing systems. Wav2Lip [2] masks the lower face of a video and inpaints it conditioned on audio, achieving lip sync close to real recordings. VideoReTalking [3] adds expression normalization and face enhancement on top of a similar approach, producing cleaner results. SadTalker [4] goes further, predicting 3D morphable model coefficients from audio to get more natural head movement than purely pixel-level methods. All three produce good output, but they assume pre-recorded video and offline compute. None are designed for a live call.

The parametric side of the literature is more relevant to Mime’s approach. VOCA [5] showed you can map audio features directly to a structured face representation, specifically 3D mesh vertex displacements, and animate a wide range of identities from a single model. Mime follows the same principle but targets ARKit blendshape coefficients, which are lighter to infer and render directly on avatar-based output. For the speech side, Whisper [6] provides a practical multilingual transcription foundation, trained on 680,000 hours of audio across dozens of languages with near-human accuracy in a zero-shot setting.

The gap none of these works address is the *full pipeline*: translation, synthesis, and synchronized facial animation running together in real time. That is the specific scope Mime targets.

Chapter 2

Methods

2.1 Methods from Neural Networks

Mime is implemented as two coupled neural pipelines: a **speech translation pipeline** (STT $\rightarrow$ MT $\rightarrow$ TTS) and an **audio-to-blendshape (ABS) pipeline** that drives a 3D ARKit avatar in real time.

2.1.1 Speech Translation Pipeline

Speech-to-Text (STT): The microphone stream is captured at 16 kHz mono and segmented online using an RMS-based voice activity strategy. Chunks are finalized after a silence window and transcribed with Groq’s `whisper-large-v3-turbo`. This chunked design reduces unnecessary API calls and keeps inference focused on complete utterances.

Machine Translation (MT): Each English utterance is translated to conversational French using `llama-3.3-70b-versatile` through a constrained system prompt. A short rolling context buffer is included to improve coherence for fragments and pronoun resolution during live dialogue.

Text-to-Speech (TTS): French text is synthesized by Inworld TTS (default model `inworld-tts-1.5-mini`, French voice). Audio is streamed as LINEAR16 at 48 kHz so playback can begin before sentence completion, which lowers perceived latency in meetings.

2.1.2 Audio-to-Blendshape Neural Model

The visual pipeline predicts 52 ARKit blendshape coefficients from audio. The training notebook uses a strided 1D convolutional frontend for local phonetic feature extraction, followed by:

1. a temporal depthwise mixing block,
2. sinusoidal positional encoding,
3. a Transformer encoder for long-range temporal context,
4. a regression head with sigmoid outputs to constrain coefficients to $[0, 1]$.

Given predicted sequence $\hat{Y}$ and target sequence Y , the core objective is *masked mean squared error* over valid (unpadded) frames:

$$\mathcal{L}_{\text{mse}} = \frac{1}{N} \sum_{t \in \Omega} \|\hat{Y}_t - Y_t\|_2^2$$

To encourage realistic temporal dynamics and reduce jitter, training adds a velocity consistency term:

$$\mathcal{L}_{\text{vel}} = \frac{1}{N'} \sum_{t \in \Omega'} \|(\hat{Y}_t - \hat{Y}_{t-1}) - (Y_t - Y_{t-1})\|_2^2$$

with total loss:

$$\mathcal{L} = \mathcal{L}_{\text{mse}} + \lambda \mathcal{L}_{\text{vel}}, \quad \lambda = 0.35$$

Training uses AdamW, OneCycleLR scheduling, mixed precision on CUDA, gradient clipping, and early stopping.

2.1.3 Real-Time Inference and Rendering

At runtime, the client uses an optimized ABS inference path with a sliding audio window and queue-based frame updates. Predicted blendshape vectors are remapped from BEAT ordering to canonical ARKit ordering, then applied to morph targets in a rigged GLB avatar via an offscreen renderer.

A transport bridge publishes:

1. rendered avatar frames to a virtual camera, and
2. synthesized TTS audio to a virtual cable device,

allowing direct integration with Zoom-compatible conferencing workflows.

2.2 Dataset and Data Processing

2.2.1 Primary Dataset

The project is trained on the **BEAT multimodal talking dataset** [7], using locally downloaded paired audio and blendshape annotations. In the current experiments, the active split comes from the English subset (`beat_english_v0.2.1`) available under the repository data directory.

2.2.2 Preprocessing Pipeline

Each sample is formed from a side-by-side `.wav` and `.json` pair:

-
1. Audio is loaded as mono and resampled to 16 kHz.
 2. Per-sample clip duration is capped (10 seconds in the current notebook setup) for stable training throughput.
 3. JSON frame weights are parsed and normalized to a fixed 52-dimensional target vector per frame (pad/truncate when needed).
 4. Targets are aligned to the effective audio duration using 60 FPS frame logic.
 5. Variable-length sequences are padded per batch with explicit length tensors for masked losses.

The training pipeline currently uses an 80/20 train/validation split with a fixed random seed for reproducibility.

2.3 Validation Strategy and Experiments

2.3.1 Model-Level Validation

Validation is performed every epoch with sequence masking to avoid padding bias. The primary tracked metrics are:

1. **Masked MSE**: main optimization-aligned error metric.
2. **Masked MAE**: interpretable average absolute coefficient error.
3. **Threshold Accuracy**: fraction of coefficients within an absolute error threshold (default 0.05).

TensorBoard logging is used to track train/validation loss, accuracy, MAE, learning rate, and trial summaries. Best checkpoints are selected by validation loss with early stopping patience and minimum-delta logic.

2.3.2 Hyperparameter Experiments

The notebook supports structured trial sweeps over learning rate, hidden dimension, Transformer depth/heads, dropout, batch size, and epoch count. Each trial writes isolated logs and checkpoints, enabling direct comparison of convergence behavior and best-epoch performance.

2.3.3 System-Level Validation

Beyond offline metrics, the system is validated in *end-to-end live runs*:

1. microphone speech $\rightarrow$ transcription,

-
2. English-to-French translation,
 3. streamed French synthesis,
 4. synchronized avatar mouth motion and Zoom bridge output.

Operational logs (for example first audio chunk latency, per-epoch ETA, and module-level status) are used to assess practical real-time behavior and robustness under device reconnect or stream interruptions.

Chapter 3

Results

3.1 Quantitative Data

All of the following results can be found in the results folder of the project GitHub repository, if too small to view here: <https://github.com/FinityFly/tobio/tree/main/results>.

3.2 Qualitative Results

Chapter 4

Discussion

4.1 Limitations and Challenges

Tobio's development came with its fair share of challenges. The most significant was achieving an accurate court segmentation system, finding a method for consistent player identification, and reliable event detection.

While the player tracking model is good at finding players (mAP50 of 0.94), the bounding boxes aren't always precise (mAP50-95 of 0.61) and the model often swapped player IDs. In the chaos of a real game, similar-looking players sometimes led to devastating identity switches, and would have to be manually corrected in post.

Moreover, the models are also only as good as their training data. The court segmentation dataset, for example, was fairly clean and didn't include the messy, overlapping lines of multi-use gyms where many amateur teams play. This means performance could dip in those more common, imperfect environments.

In addition, compiling the 3D spatial coordinates of the ball was one of the biggest initial constraints, since our objective was to use a single camera viewpoint. Though exact XYZ could be obtained, the depth information was unreliable without more powerful models and more research.

The most difficult challenge had to be event detection. Volleyball actions are continuous and fluid, making it hard to define clear start and end points for each event, especially in a sport as complicated and fast-paced as volleyball. The model often confused similar actions, like spikes and serves, leading to misclassifications that would need manual review. The overall system I developed was only a prototype and had some questionable logic and reliability issues due to its complexity, but I know that a well-polished system is definitely within arms reach with more time and resources.

4.2 Directions for Future Work

I definitely see a lot of potential in improving Tobio further, both in terms of turning it into a fully-polished product by improving the technical components, as well as expanding into other sports and use-cases. First of all, I'd like to improve the recall of the most important subcomponents: ball tracking, player recognition, and event detection. These form the backbone of what's possible with Tobio, so improving their reliability is a priority. Additionally, I'd like to improve my application's practicality with real competitive teams' use-cases in mind. This includes building more meaningful visualizations, more powerful and flexible user interface that allows for each teams' customizations, and curated reports for both coaches and players.

Although, other more popular sports already have implementations of similar systems, these popular sports often have paid subscriptions to such services and the non-popular sports don't receive enough attention for innovation. Volleyball is just the start line, and I'd like to examine what possibilities exist in other domains.

4.3 Implications and Impact

I believe that Tobio has the potential to democratize sports performance analysis, since one of my key morals is making it accessible to anyone, including teams that lack the resources of professional organizations. Through automating the entire process of video analysis, teams can save opportunity costs on tedious manual data collection and focus more on the other significant aspects one would much rather be doing.

4.4 Conclusions

In conclusion, this project has successfully developed a proof-of-concept AI-powered platform for volleyball video analysis. While there are still many shortcomings due to the limited scope and time constraints of this project, Tobio's end-to-end system is still able to show a glimpse into what is possible when combining computer vision, machine learning, and sports analytics. With proper development and time, Tobio has the potential to become a valuable tool for amateur and collegiate volleyball teams, all while paving the way for similar systems in other sports.

References

- [1] H. McGurk and J. MacDonald, “Hearing lips and seeing voices,” *Nature*, vol. 264, no. 5588, pp. 746–748, 1976.
- [2] K. R. Prajwal, R. Mukhopadhyay, V. P. Namboodiri, and C. Jawahar, “A lip sync expert is all you need for speech to lip generation in the wild,” in *Proceedings of the 28th ACM International Conference on Multimedia*, MM ’20, (New York, NY, USA), pp. 484–492, Association for Computing Machinery, 2020.
- [3] K. Cheng, X. Cun, Y. Zhang, M. Xia, F. Yin, M. Zhu, X. Wang, J. Wang, and N. Wang, “Videoretalking: Audio-based lip synchronization for talking head video editing in the wild,” 2022.
- [4] W. Zhang, X. Cun, X. Wang, Y. Zhang, X. Shen, Y. Guo, Y. Shan, and F. Wang, “Sadtalker: Learning realistic 3D motion coefficients for stylized audio-driven single image talking face animation,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 8652–8661, 2023.
- [5] D. Cudeiro, T. Bolkart, C. Laidlaw, A. Ranjan, and M. J. Black, “Capture, learning, and synthesis of 3D speaking styles,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 10101–10111, 2019.
- [6] A. Radford, J. W. Kim, T. Xu, G. Brockman, C. McLeavey, and I. Sutskever, “Robust speech recognition via large-scale weak supervision,” in *Proceedings of the 40th International Conference on Machine Learning*, vol. 202 of *Proceedings of Machine Learning Research*, pp. 28492–28518, PMLR, 2023.
- [7] H. Liu, Z. Zhu, N. Iwamoto, Y. Peng, Z. Li, Y. Zhou, E. Bozkurt, and B. Zheng, “Beat: A large-scale semantic and emotional multi-modal dataset for conversational gestures synthesis,” 2022.